

Intro into

facto

Agenda

- What is ecto
- Basic usage
- Protips

Agenda

- What is ecto
- Basic usage
- Protips

Ecto is a domain specific
language for writing queries
and interacting with
databases in Elixir.

Supports various databases

- PostgreSQL
- MySQL
- MSSQL
- SQLite3
- MongoDB

Ecto's guts

- **Repo**
- **Query**
- **Migrations**
- **Changesets**
- **Adapters**

Agenda

- What is ecto
- Basic usage
- Protips

Repo

- Defines a repository used by the app.
- Connects the adapter + config for the database.
- Interface for managing data in the database that delegates to `Repo.Queryable`, `Repo.Query`, and a few others.

Repo setup

```
config :lost_legends, LostLegends.Repo,  
  adapter: Ecto.Adapters.Postgres,  
  url: "postgres://garrett:garrett@localhost/lost_legends_dev",  
  pool_size: 10  
  
defmodule LostLegends.Repo do  
  use Ecto.Repo, otp_app: :lost_legends  
end
```

What do we get?

`mix ecto.create`

`mix ecto.drop`

`mix ecto.gen.migration`

`mix ecto.gen.repo`

`mix ecto.migrate`

`mix ecto.rollback`

`# Create the storage for the repo`

`# Drop the storage for the repo`

`# Generate a new migration for the repo`

`# Generate a new repository`

`# Run migrations up on a repo`

`# Rollback migrations from a repo`

Migrations

```
$ mix ecto.gen.migration create_monsters
* creating priv/repo/migrations/20160112072054_create_monsters.exs
```

```
defmodule LostLegends.Repo.Migrations.CreateMonsters do
  use Ecto.Migration

  def change do
    create table(:monsters) do
      add :name, :string
      add :desc, :text
      add :health, :integer

      timestamps
    end
  end
end
```

Schema

```
defmodule LostLegends.Monster do
  use Ecto.Schema

  schema "monsters" do
    field :name, :string
    field :desc, :string
    field :health, :integer

    timestamps
  end
end
```

Lets use this sucker

```
iex> alias LostLegends.{Repo, Monster}
iex> Repo.all Monster
[debug] SELECT mo."id", mo."name", mo."desc", mo."health", mo."inserted_at", mo."updated_at"
      FROM "monsters" AS mo [] OK query=2.7ms
[  

  %LostLegends.Monster{__meta__: #Ecto.Schema.Metadata<:loaded>,
    desc: "Drugs are bad mkay.", health: 10, id: 1,
    inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>, name: "Crackie Monster",
    updated_at: #Ecto.DateTime<2016-01-07T11:20:54Z>},
  %LostLegends.Monster{__meta__: #Ecto.Schema.Metadata<:loaded>,
    desc: "Today totally isn't Sunday, trust me.", health: 3000, id: 2,
    inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>, name: "That 12th beer",
    updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}]  

```



1100011010001101010110100100111
10001101011000100000010000101
10101100111000101000010001001
00011110000100100111101010101
01001001000111100001101001001
10001001101100111110010011000
10100110100111001110110110001
010000100000010001101100001
10101100100011010111011110110
1100001011000000110001111001
10001001110010100100010011100

Such Data

Repo.all(Model)

- Repo is the interface for the database
- No magic functions on Model

Repo Interface

Repo.all

Repo.one

Repo.get

Repo.insert

Repo.update(_all)

Repo.delete(_all)

Repo.transaction

Repo.preload

and more...!



The thing **Repo** takes

Query

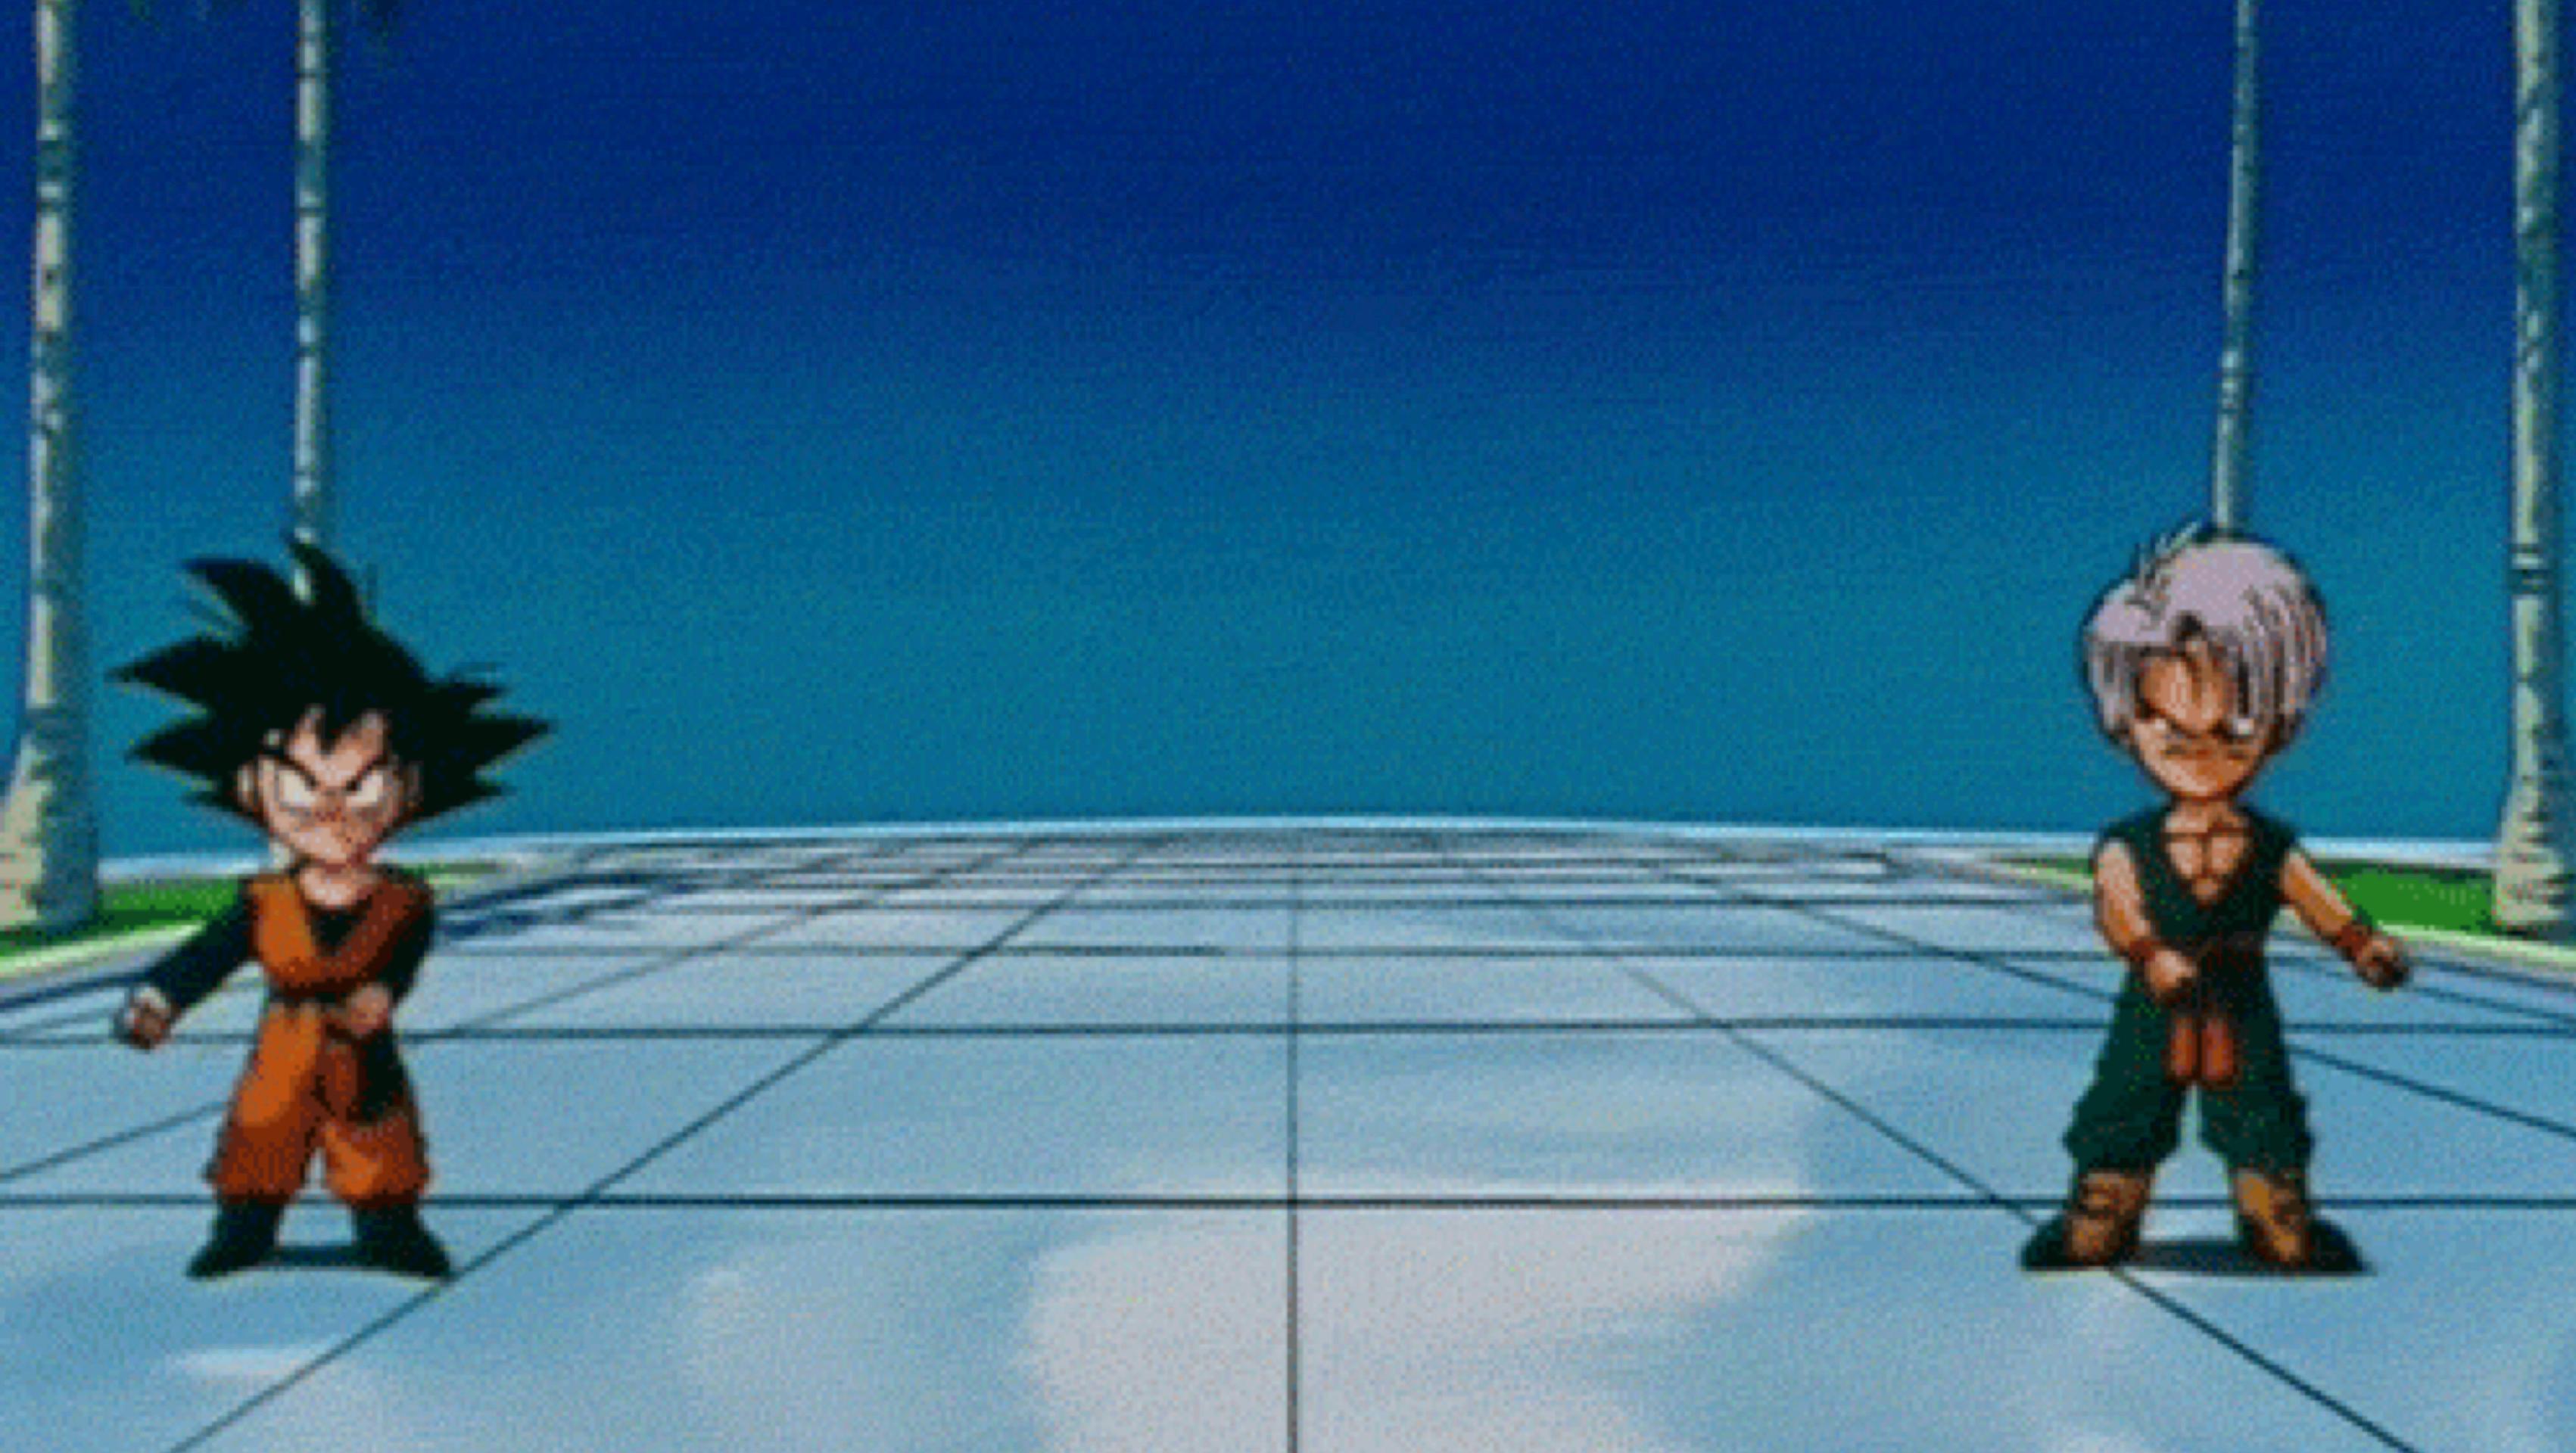
```
import Ecto.Query, only: [from: 2]

query = from m in Monster, where: m.health > 10
```

```
Repo.all(query)
```

composition

When two become one



composition

```
import Ecto.Query, only: [from: 2]
```

```
such_health_query = from m in Monster, where: m.health > 10
```

```
recent_ms_query = from m in such_health_query, # notice me  
                   where: m.inserted_at > datetime_add(^Ecto.DateTime.utc, -1, "month")
```

```
Repo.all(recent_ms_query)
```

Real example

```
def required_assignments(date, remote_id) do
  filter_by_status(required_status)
  |> for_student(remote_id)
  |> current(date)
  |> Dogfood.Repo.all
end

defp required_status, do: "required"
defp filter_by_status(status) do
  from assignment in Dogfood.Assignment,
    where: assignment.status == ^status
end

defp for_student(query, remote_id) do
  from assignment in query,
    where: assignment.student_id == ^remote_id # RE: student_id == MS: remote_id
end

defp current(query, %Ecto.DateTime{} = date) do
  from assignment in query,
    where: assignment.start_date <= ^date and ^date < assignment.end_date
end
```

Preloading Records

Preventing your looping problems since 2013

Preloading Records

```
iex> battle = Repo.get Battle, 1
[debug] SELECT bo."id", bo."monster_id", bo."inserted_at", bo."updated_at"
          FROM "battles" AS bo WHERE (bo."id" = $1) [1] OK query=0.7ms
%LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 1,
  inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>,
  monster: #Ecto.Association.NotLoaded<association :monster is not loaded>,
  monster_id: 1, updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}
iex> battle.monster
#Ecto.Association.NotLoaded<association :monster is not loaded>
```

Preloading Records

```
iex> battle = Repo.get(Battle, 1) |> Repo.preload(:monster)

[debug] SELECT bo."id", bo."monster_id", bo."inserted_at", bo."updated_at"
      FROM "battles" AS bo WHERE (bo."id" = $1) [1] OK query=0.7ms
[debug] SELECT mo."id", mo."name", mo."desc", mo."health", mo."inserted_at", mo."updated_at"
      FROM "monsters" AS mo WHERE (mo."id" IN ($1)) [1] OK query=0.5ms
%LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 1,
inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>,
monster: %LostLegends.Monster{__meta__: #Ecto.Schema.Metadata<:loaded>,
  desc: "Drugs are bad mkay.", health: -9, id: 1,
  inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>, name: "Crackie Monster",
  updated_at: #Ecto.DateTime<2016-01-07T11:20:54Z>}, monster_id: 1,
updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}

iex> battle.monster
%LostLegends.Monster{__meta__: #Ecto.Schema.Metadata<:loaded>,
desc: "Drugs are bad mkay.", health: 10, id: 1,
inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>, name: "Crackie Monster",
updated_at: #Ecto.DateTime<2016-01-07T11:20:54Z>}
```

```
    iex> battle = Repo.get(Battle, 1) |> Repo.preload(monster: :battles)

[debug] SELECT bo."id", bo."monster_id", bo."inserted_at", bo."updated_at"
      FROM "battles" AS bo WHERE (bo."id" = $1) [1] OK query=0.6ms
[debug] SELECT mo."id", mo."name", mo."desc", mo."health", mo."inserted_at", mo."updated_at"
      FROM "monsters" AS mo WHERE (mo."id" IN ($1)) [1] OK query=0.3ms queue=0.1ms
[debug] SELECT bo."id", bo."monster_id", bo."inserted_at", bo."updated_at"
      FROM "battles" AS bo WHERE (bo."monster_id" IN ($1)) ORDER BY bo."monster_id" [1] OK query=0.4ms
%LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 1,
  monster: %LostLegends.Monster{__meta__: #Ecto.Schema.Metadata<:loaded>,
    battles: [%LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 1,
      monster: #Ecto.Association.NotLoaded<association :monster is not loaded>,
      monster_id: 1, updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>},
      %LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 2,
        monster: #Ecto.Association.NotLoaded<association :monster is not loaded>,
        monster_id: 1, updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}],
    desc: "Drugs are bad mkay.", health: -9, id: 1,
    inserted_at: #Ecto.DateTime<2016-01-07T11:02:43Z>, name: "Crackie Monster",
    updated_at: #Ecto.DateTime<2016-01-07T11:20:54Z>}, monster_id: 1,
    updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}

    iex> battle.monster.battles
[%LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 1,
  monster: #Ecto.Association.NotLoaded<association :monster is not loaded>,
  monster_id: 1, updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>},
  %LostLegends.Battle{__meta__: #Ecto.Schema.Metadata<:loaded>, id: 2,
    monster: #Ecto.Association.NotLoaded<association :monster is not loaded>,
    monster_id: 1, updated_at: #Ecto.DateTime<2016-01-07T11:02:43Z>}]
```

Agenda

- What is ecto
- Basic usage
- Protips

Protips

Fragments

```
def unpublished_by_title(title) do
  from p in Post,
  where: is_nil(p.published_at) and
    fragment("downcase(?)", p.title) == ^title
end
```

Protips

Piping

```
# no worky
from u in User, where: u.admin == true
| > Repo.all
```

```
# worky
query = from u in User, where: u.admin == true
query |> Repo.all
```

Protips

no HABTM, but HMT works fine

```
schema "recurrings" do
  field :frequency, :string
  has_many :recurring_tags, BudgetApi.RecurringTag
  has_many :tags, through: [:recurring_tags, :tag]
end

schema "transactions" do
  field :amount, :float, default: 0.0
  has_many :transaction_tags, BudgetApi.TransactionTag
  has_many :tags, through: [:transaction_tags, :tag]
end

schema "recurrings_tags" do
  belongs_to :recurring, BudgetApi.Recurring, references: :id
  belongs_to :tag, BudgetApi.Tag, references: :id
end

schema "transactions_tags" do
  belongs_to :transaction, BudgetApi.Transaction, references: :id
  belongs_to :tag, BudgetApi.Tag, references: :id
end

schema "tags" do
  field :tag, :string
  has_many :recurring_tags, BudgetApi.RecurringTag
  has_many :recurrings, through: [:recurring_tags, :recurring]
  has_many :transaction_tags, BudgetApi.TransactionTag
  has_many :transactions, through: [:transaction_tags, :transaction]
end
```

Protips

"`Repo.first` and `Repo.last` have been historically confusing on Rails because many databases do not give a guarantee of ordering. [...]

That's why Ecto provides `Repo.one` that does not give any idea of ordering."

— José Valim

```
def first(query, opts \\ []) do
  Repo.one(from(q in query, order_by: [asc: q.id], limit: 1), opts)
end
```

```
def last(query, opts \\ []) do
  Repo.one(from(q in query, order_by: [desc: q.id], limit: 1), opts)
end
```

Protips

jsonb columns

```
create table(:assignments) do
  add :assignment_details, :map
end

schema "assignments" do
  field :assignment_details, :map
end
```

```
iex> {:ok, assignment} = Repo.insert(%Assignment{assignment_details: %{awesome: "elixir"}})
iex> assignment.assignment_details["awesome"] # "elixir"
```

thanks

@gogogarrett