

From Fast to Ludicrous

Performance Tuning in Elixir

Josh Price

Elixir is Fast

When it isn't Fast enough



Ludicrous Speed!

Performance Tuning

**If it's fast enough
don't worry about it
(really)**

The real problem is that programmers have spent far too much time worrying about efficiency in the wrong places and at the wrong times; premature optimization is the root of all evil (or at least most of it) in programming.

-- Donald Knuth

Class Solving

- Need to allocate students into classes with teachers optimally
- Typically done by hand slowly (takes weeks)
- Teacher and Parent pair/separation requests
- Ensure gender balance and custom characteristics
- At least 1 friend in class
- NP-Hard (don't need perfect solution though)

Some Background

- Solving classes for students is very CPU heavy
- Increased server costs
- Low max number of concurrent solves
- Synchronous: Don't want to keep user waiting too long
- 16Gb 8 core MBP is faster than a 2 core VM with 2Gb RAM

Finished in 9.1 seconds

125 tests, 0 failures

Possible Problem Areas

- Wrong data structure
- Wrong algorithm
- Doing the same thing more than once
- Throwing away results you could use again
- All of the above

Strategies

1. ~Change the whole algorithm (not yet too risky)~
2. Find the slow bits by profiling
3. Isolate and optimise
4. Measure improvements
5. Test (make sure it works!)
6. Repeat 2-5 until fast enough

Profiling code

- Needed for finding the bottlenecks in your code
- Where does the program spend it's time?
- Slow functions
- Fast functions called many, many times
- Use the 80/20 rule
- Make sure it's the slowest improvable code
- Don't fix the unimportant stuff

Elixir/Erlang Profilers

- cprof
- eprof
- fprof

cprof

- Nested function call counts across whole program
- Useful to see what functions are called a lot
- Good overview of your modules
- No timing info

mix profile.cprof

```
mix profile.cprof -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

	CNT	
Total	9	
Enum	6	<--
Enum."-map/2-lists^map/1-0-"/2	4	
Enum.reverse/1	1	
Enum.map/2	1	
:erlang	2	<--
:erlang.trace_pattern/3	2	
:elixir_compiler_3	1	<--
anonymous fn/0 in :elixir_compiler_3.__FILE__/1	1	

Profile done over 16880 matching functions

--matching <Mod.fun/arity>

```
mix profile.cprof --matching Enum -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

	CNT	
Total	6	
Enum	6	<--
Enum." map /2 lists ^{^map} /1 0 "/2	4	
Enum.reverse/1	1	
Enum.map/2	1	

Profile done over 399 matching functions

eprof

- Time taken for function calls with counts
- More info than cprof
- Sorted by total time
- *Start at the bottom of the list*

mix profile.eprof

```
mix profile.eprof -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

Profile results of #PID<0.143.0>

#	CALLS	% TIME	μS/CALL	
Total	13	100.0	17	1.31
:lists.reverse/2	1	0.00	0	0.00
:erlang.make_fun/3	1	5.88	1	1.00
Enum.map/2	1	5.88	1	1.00
anonymous fn/0 in :elixir_compiler_1.__FILE__/1	1	5.88	1	1.00
:erlang.apply/2	1	11.76	2	2.00
:erlang.integer_to_binary/1	3	17.65	3	1.00
Enum."-map/2-lists^map/1-0-"/2	4	23.53	4	1.00
Enum.reverse/1	1	29.41	5	5.00

Profile done over 8 matching functions

--matching Enum

```
mix profile.eprof --matching Enum -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

Profile results of #PID<0.143.0>

#		CALLS	% TIME	μS/CALL	
Total		6	100.0	7	1.17
Enum.map/2		1	14.29	1	1.00
Enum."-map/2-lists^map/1-0-"/2		4	42.86	3	0.75
Enum.reverse/1		1	42.86	3	3.00

Profile done over 3 matching functions

fprof

- Nested Time taken for function calls with counts
- Shows accumulated time (ACC)
- Sorted by total time (ACC)
- *Start at the TOP of the list*
- Warning: high tracing cost means could take a long time
- Much more detail

mix profile.fprof

```
mix profile.fprof -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

	CNT	ACC (ms)	OWN (ms)
Total	14	0.043	0.043
:fprof.apply_start_stop/4	0	0.043	0.007
anonymous fn/0 in :elixir_compiler_1.__FILE__/1	1	0.032	0.005
Enum.map/2	1	0.020	0.002
Enum."-map/2-lists^map/1-0-"/2	4	0.018	0.015
Enum.reverse/1	1	0.005	0.004
:fprof."-apply_start_stop/4-after\$^1/0-0-"/3	1	0.004	0.004
:erlang.integer_to_binary/1	3	0.003	0.003
:erlang.make_fun/3	1	0.002	0.002
:lists.reverse/2	1	0.001	0.001
:undefined	0	0.000	0.000
:suspend	1	0.000	0.000

--callers graph

```
mix profile.fprof --callers -e "[1, 2, 3] |> Enum.reverse |> Enum.map(&Integer.to_string/1)"
```

Warmup...

	CNT	ACC (ms)	OWN (ms)	
Total	14	0.043	0.043	
anonymous fn/0 in :elixir_compiler_1.__FILE__/1	1	0.021	0.001	
Enum.map/2	1	0.021	0.001	<--
Enum."-map/2-lists^map/1-0-"/2	1	0.020	0.011	
Enum.map/2	1	0.020	0.011	
Enum."-map/2-lists^map/1-0-"/2	3	0.000	0.006	
Enum."-map/2-lists^map/1-0-"/2	4	0.020	0.017	<--
:erlang.integer_to_binary/1	3	0.003	0.003	
Enum."-map/2-lists^map/1-0-"/2	3	0.000	0.006	
anonymous fn/0 in :elixir_compiler_1.__FILE__/1	1	0.005	0.004	
Enum.reverse/1	1	0.005	0.004	<--
:lists.reverse/2	1	0.001	0.001	
Enum."-map/2-lists^map/1-0-"/2	3	0.003	0.003	
:erlang.integer_to_binary/1	3	0.003	0.003	<--
Enum.reverse/1	1	0.001	0.001	
:lists.reverse/2	1	0.001	0.001	<--

Measuring with Microbenchmarks

- Benchee <https://github.com/PragTob/benchee>

```
list = Enum.to_list(1..10_000)
map_fun = fn(i) -> [i, i * i] end
```

```
Benchee.run(%{
  "flat_map"      => fn -> Enum.flat_map(list, map_fun) end,
  "map.flatten"  => fn -> list |> Enum.map(map_fun) |> List.flatten end
}, time: 10, memory_time: 2)
```

mix run samples/run.exs

Operating System: Linux"

CPU Information: Intel(R) Core(TM) i7-4790 CPU @ 3.60GHz

Number of Available Cores: 8

Available memory: 15.61 GB

Elixir 1.6.4

Erlang 20.3

Benchmark suite executing with the following configuration:

warmup: 2 s

time: 10 s

memory time: 2 s

parallel: 1

inputs: none specified

Estimated total run time: 28 s

mix run samples/run.exs

Benchmarking flat_map...
Benchmarking map.flatten...

Name	ips	average	deviation	median	99th %
flat_map	2.31 K	433.25 μ s	\pm 8.64%	428 μ s	729 μ s
map.flatten	1.22 K	822.22 μ s	\pm 16.43%	787 μ s	1203 μ s

Comparison:

flat_map	2.31 K
map.flatten	1.22 K – 1.90x slower

Memory usage statistics:

Name	Memory usage
flat_map	625.54 KB
map.flatten	781.85 KB – 1.25x memory usage

Reminders

- Only optimize as last resort
- Good tests stop you breaking things
- Only work on the slowest code
- Measure!

Thanks!